

RED WINE QUALITY ANALYSIS

In [4]:

```
import pandas as pd
```

In [5]:

```
data = pd.read_csv("D:\Data Analytics Project\Red Wine Quality\winequality-r
```

In [7]:

```
data.head(10)
```

Out[7]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	1
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	1

In [8]:

```
data.tail(10)
```

Out[8]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
1589	6.6	0.725	0.20	7.8	0.073	29.0	79.0	0.99770	3.29	0.54	
1590	6.3	0.550	0.15	1.8	0.077	26.0	35.0	0.99314	3.32	0.82	
1591	5.4	0.740	0.09	1.7	0.089	16.0	26.0	0.99402	3.67	0.56	
1592	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1593	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	

DATA PREPROCESSING

```
In [10]: print('    Size of the table')
print('No.of Rows:',data.shape[0])
print('No.of Columns:',data.shape[1])
```

```
    Size of the table
No.of Rows: 1599
No.of Columns: 12
```

```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   fixed acidity         1599 non-null   float64
 1   volatile acidity      1599 non-null   float64
 2   citric acid           1599 non-null   float64
 3   residual sugar        1599 non-null   float64
 4   chlorides             1599 non-null   float64
 5   free sulfur dioxide   1599 non-null   float64
 6   total sulfur dioxide  1599 non-null   float64
 7   density              1599 non-null   float64
 8   pH                   1599 non-null   float64
 9   sulphates            1599 non-null   float64
10   alcohol              1599 non-null   float64
11   quality              1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [12]: data.describe()
```

```
Out[12]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total di
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.0
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.4
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.8
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.0
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.0
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.0
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.0
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.0

```
In [13]: # Checking Null values  
data.isnull().sum()
```

```
Out[13]: fixed acidity          0  
volatile acidity       0  
citric acid            0  
residual sugar         0  
chlorides              0  
free sulfur dioxide    0  
total sulfur dioxide   0  
density                0  
pH                    0  
sulphates              0  
alcohol                0  
quality                0  
dtype: int64
```

```
In [15]: quality = data['quality'].value_counts()  
quality
```

```
Out[15]: 5    681  
        6    638  
        7    199  
        4     53  
        8     18  
        3     10  
Name: quality, dtype: int64
```


DATA VISUALIZATION AND EDA

```
In [30]: # Quality Distribution
quality = data.groupby('quality').size()
import matplotlib.pyplot as plt

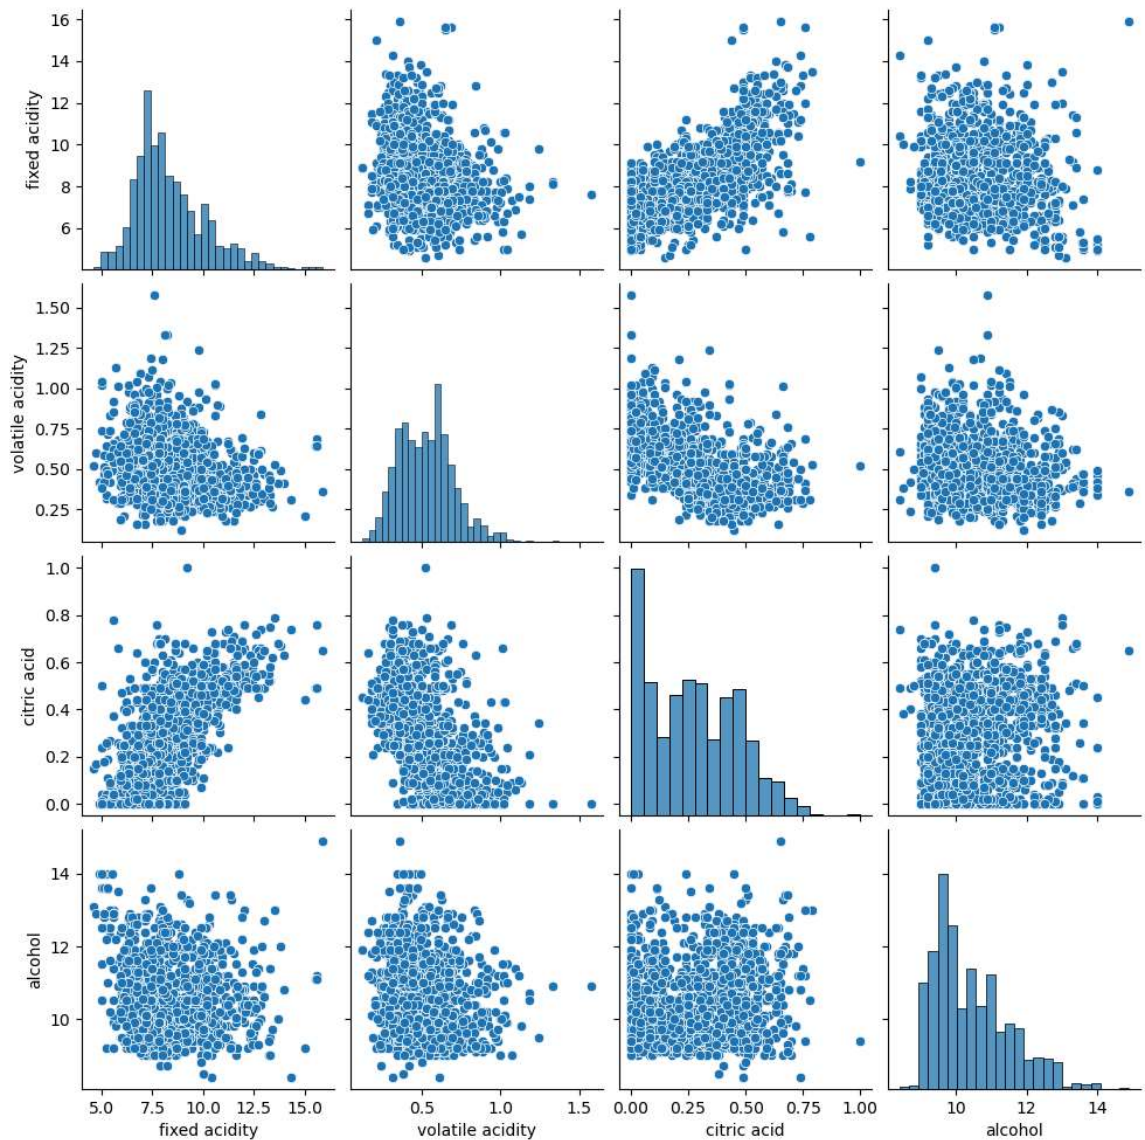
plt.figure(figsize=(10,6))
quality.plot.bar(color=['blue','red','green','salmon','cyan','purple'],edgecolor='black')
plt.title('Quality Distribution',fontsize=18,fontweight='bold',color='darkred')
plt.xlabel('Quality',fontsize=12,fontweight='bold',color='darkblue')
plt.ylabel('No. of Counts',fontsize=12,fontweight='bold',color='darkblue')
plt.tight_layout()
plt.show()
```



```
In [47]: # Distribution of pH Level
import seaborn as sns
plt.figure(figsize=(8,6))
sns.histplot(data['pH'],kde=True,bins=20)
plt.title('Distribution of pH Levels in Wine',fontsize=18,fontweight='bold')
plt.xlabel('pH Level',fontsize=14,color='darksalmon',fontweight='bold')
plt.ylabel('Frequency',fontsize=14,color='darksalmon',fontweight='bold')
plt.show()
```




```
In [51]: # Relationship between key features
sns.pairplot(data[['fixed acidity', 'volatile acidity', 'citric acid', 'alcohol']])
plt.show()
```



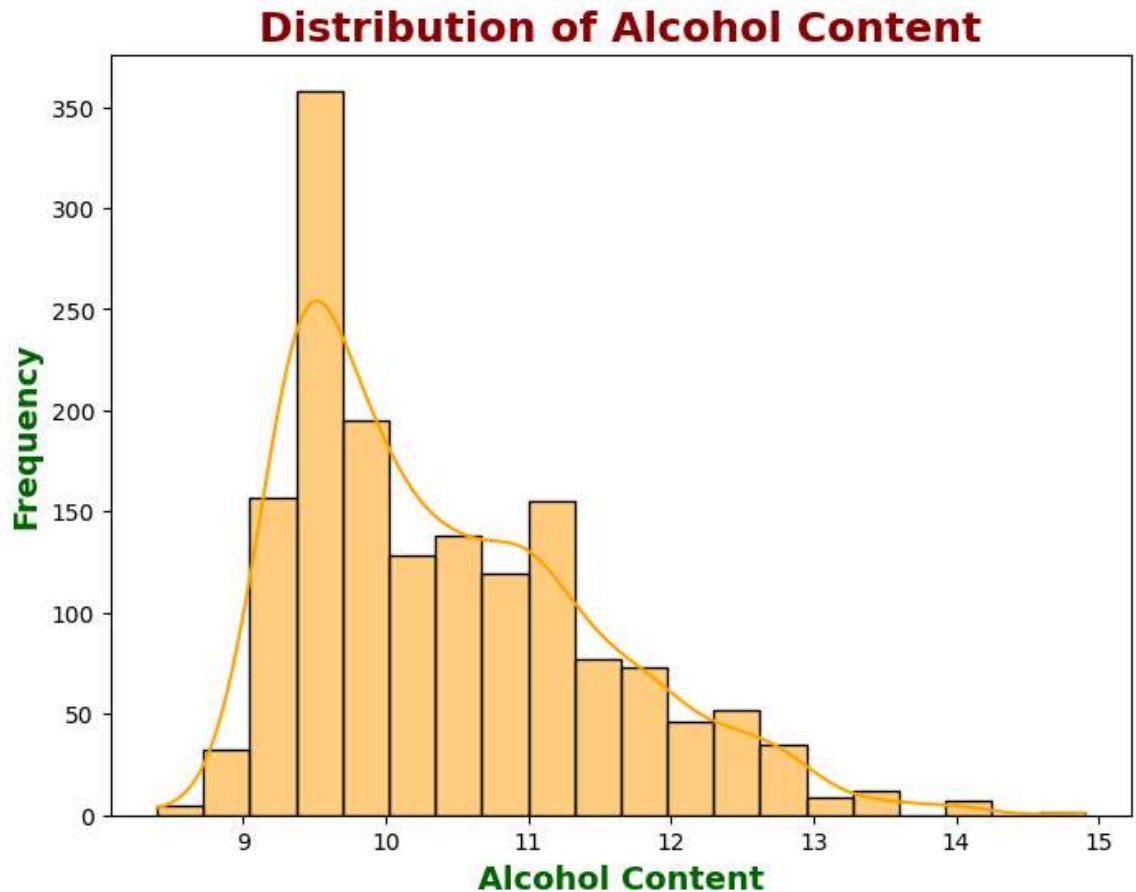
```
In [53]: # Correlation Matrix
corr_matrix = data.corr()
plt.figure(figsize=(12,8))
sns.heatmap(corr_matrix,annot=True,cmap='coolwarm',linewidth=0.5)
plt.title('Correlation Matrix of Features')
plt.show()
```

C:\Users\ROHITH DP\AppData\Local\Temp\ipykernel_7932\2098915487.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = data.corr()
```



```
In [56]: # Alcohol Distribution
plt.figure(figsize=(8,6))
sns.histplot(data['alcohol'],kde=True,bins=20,color='orange')
plt.title('Distribution of Alcohol Content',fontsize=18,
          fontweight='bold',color='darkred')
plt.xlabel('Alcohol Content',fontsize=14,fontweight='bold',color='darkgreen')
plt.ylabel('Frequency',fontsize=14,fontweight='bold',color='darkgreen')
plt.show()
```



CLASSIFICATION

```
In [58]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_
```



```
In [37]: bins = [2,4,6,8]
labels = ['low','medium','high']
data['quality_category'] = pd.cut(data['quality'],bins=bins,labels=labels,include_lowest=True)
data = data.drop('quality',axis=1)

# Separate features and target variables
x = data.drop('quality_category',axis=1)
y = data['quality_category']

# Split, Train and Test
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.2,random_state=42)
```

Random Forest Classifier

```
In [67]: rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
y_pred = rfc.predict(x_test)
rfc_accuracy = accuracy_score(y_test,y_pred)
rfc_report = classification_report(y_test,y_pred)
print("Accuracy:",rfc_accuracy)
print("Classification Report:")
print(rfc_report)

rfc_conf_matrix = confusion_matrix(y_test,y_pred)
print("Confusion Matrix:\n",rfc_conf_matrix)
plt.figure(figsize=(10,6))
sns.heatmap(rfc_conf_matrix,annot=True,fmt='d',cmap='Reds')
plt.title('Confusion Matrix',fontsize=18,fontweight='bold',color='teal')
plt.xlabel('<-----Predicted----->',fontsize=14,color='darkgreen')
plt.ylabel('<-----Actual----->',fontsize=14,color='darkgreen')
plt.show()
```

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))


Accuracy: 0.859375

Classification Report:

	precision	recall	f1-score	support
high	0.66	0.57	0.61	47
low	0.00	0.00	0.00	11
medium	0.89	0.95	0.92	262
accuracy			0.86	320
macro avg	0.52	0.51	0.51	320
weighted avg	0.82	0.86	0.84	320

Confusion Matrix:

```
[[ 27  0 20]
 [  0  0 11]
 [ 14  0 248]]
```



Decision Tree Classifier

```
In [69]: dtc = DecisionTreeClassifier()
dtc.fit(x_train,y_train)
y_pred = dtc.predict(x_test)
dtc_accuracy = accuracy_score(y_test,y_pred)
dtc_report = classification_report(y_test,y_pred)
print('Accuracy:',dtc_accuracy)
print("Classification Report:")
print(dtc_report)
dtc_conf_matrix = confusion_matrix(y_test,y_pred)
print('Confusion Matrix:\n',dtc_conf_matrix)
plt.figure(figsize=(10,6))
sns.heatmap(dtc_conf_matrix,annot=True,fmt='d',cmap='Greens')
plt.title('Confusion Matrix',fontsize=18,color='darkred')
plt.xlabel('<-----Predicted----->',fontsize=14,color='darkgreen')
plt.ylabel('<-----Actual----->',fontsize=14,color='darkgreen')
plt.show()
```


Accuracy: 0.796875

Classification Report:

	precision	recall	f1-score	support
high	0.51	0.57	0.54	47
low	0.10	0.09	0.10	11
medium	0.88	0.87	0.87	262
accuracy			0.80	320
macro avg	0.50	0.51	0.50	320
weighted avg	0.80	0.80	0.80	320

Confusion Matrix:

```
[[ 27  0 20]
 [  0  1 10]
 [ 26  9 227]]
```



Logistic Regression

```
In [100]: lr = LogisticRegression()
lr.fit(x_train,y_train)
y_pred = lr.predict(x_test)
lr_accuracy = accuracy_score(y_test,y_pred)
lr_report = classification_report(y_test,y_pred)
print("Accuracy:",lr_accuracy)
print("Classification Report")
print(lr_report)
lr_conf_matrix = confusion_matrix(y_test,y_pred)
print('Confusion Matrix:\n',lr_conf_matrix)
plt.figure(figsize=(10,6))
sns.heatmap(lr_conf_matrix,annot=True,fmt='d',cmap='Blues')
plt.title('Confusion Matrix',fontsize=18,fontweight='bold',color='darkred')
plt.xlabel('<-----Predicted----->',fontsize=14,color='purple')
plt.ylabel('<-----Actual----->',fontsize=14,color='purple')
plt.show()
```

E:\Anaconda Software\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

E:\Anaconda Software\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```


Accuracy: 0.809375

Classification Report

	precision	recall	f1-score	support
high	0.45	0.28	0.34	47
low	0.00	0.00	0.00	11
medium	0.85	0.94	0.89	262
accuracy			0.81	320
macro avg	0.43	0.41	0.41	320
weighted avg	0.76	0.81	0.78	320

Confusion Matrix:

```
[[ 13  0 34]
 [  0  0 11]
 [ 16  0 246]]
```



```
In [99]: plt.figure()  
plt.text(0.5,0.6,'Thank You',fontsize=30,ha='center',va='center')  
plt.text(0.5,0.4,'CognoRise Infotech',fontsize=40,fontweight='bold',color='r')  
plt.axis('off')  
plt.show()
```

Thank You

CognoRise Infotech